

## ▼ Using more sophisticated images with Convolutional Neural

In the previous lesson you saw how to use a CNN to make your recognition of the handwriting dig to the next level, recognizing real images of Cats and Dogs in order to classify an incoming image handwriting recognition made your life a little easier by having all the images be the same size and shape. Real-world images aren't like that -- they're in different shapes, aspect ratios etc, and they're usually not square. So, as part of the task you need to process your data -- not least resizing it to be uniform in shape. You'll follow these steps:


1. Explore the Example Data of Cats and Dogs
2. Build and Train a Neural Network to recognize the difference between the two
3. Evaluate the Training and Validation accuracy

## ▼ Explore the Example Data

Let's start by downloading our example data, a .zip of 2,000 JPG pictures of cats and dogs, and ex

**NOTE:** The 2,000 images used in this exercise are excerpted from the ["Dogs vs. Cats" dataset](#) available on Kaggle. Here, we use a subset of the full dataset to decrease training time for educational purposes.

```
1  !wget --no-check-certificate \
2    https://storage.googleapis.com/mledu-datasets/cats_and_dogs_filtered.zip \
3    -O /tmp/cats_and_dogs_filtered.zip
```

 --2020-03-15 18:41:48-- [https://storage.googleapis.com/mledu-datasets/cats\\_and\\_dogs\\_filtered.zip](https://storage.googleapis.com/mledu-datasets/cats_and_dogs_filtered.zip)  
 Resolving storage.googleapis.com (storage.googleapis.com)... 74.125.20.128, 209.85.229.104  
 Connecting to storage.googleapis.com (storage.googleapis.com)|74.125.20.128|:443...  
 HTTP request sent, awaiting response... 200 OK  
 Length: 68606236 (65M) [application/zip]  
 Saving to: '/tmp/cats\_and\_dogs\_filtered.zip'

```
/tmp/cats_and_dogs_ 100%[=====>] 65.43M 75.6MB/s in 0.9s
```

2020-03-15 18:41:49 (75.6 MB/s) - '/tmp/cats\_and\_dogs\_filtered.zip' saved [68606236]

The following python code will use the OS library to use Operating System libraries, giving you access to the file system, allowing you to unzip the data.

```
1  import os
2  import zipfile
3
4  local_zip = '/tmp/cats_and_dogs_filtered.zip'
5
6  zip_ref = zipfile.ZipFile(local_zip, 'r')
```

```

7
8  zip_ref.extractall('/tmp')
9  zip_ref.close()

```

The contents of the .zip are extracted to the base directory /tmp/cats\_and\_dogs\_filtered, which contains subdirectories for the training and validation datasets (see the [Machine Learning Crash Course](#) for details), which in turn each contain cats and dogs subdirectories.

In short: The training set is the data that is used to tell the neural network model that 'this is what' etc. The validation data set is images of cats and dogs that the neural network will not see as part of how badly it does in evaluating if an image contains a cat or a dog.

One thing to pay attention to in this sample: We do not explicitly label the images as cats or dogs. In the example earlier, we had labelled 'this is a 1', 'this is a 7' etc. Later you'll see something called an ImageGenerator to read images from subdirectories, and automatically label them from the name of that subdirectory. For example, a directory containing a 'cats' directory and a 'dogs' one. ImageGenerator will label the images appropriately.

Let's define each of these directories:

```

1  base_dir = '/tmp/cats_and_dogs_filtered'
2
3  train_dir = os.path.join(base_dir, 'train')
4  validation_dir = os.path.join(base_dir, 'validation')
5
6  # Directory with our training cat/dog pictures
7  train_cats_dir = os.path.join(train_dir, 'cats')
8  train_dogs_dir = os.path.join(train_dir, 'dogs')
9
10 # Directory with our validation cat/dog pictures
11 validation_cats_dir = os.path.join(validation_dir, 'cats')
12 validation_dogs_dir = os.path.join(validation_dir, 'dogs')
13


```

Now, let's see what the filenames look like in the cats and dogs train directories (file naming convention in each directory):

```

1  train_cat_fnames = os.listdir( train_cats_dir )
2  train_dog_fnames = os.listdir( train_dogs_dir )
3
4  print(train_cat_fnames[:10])
5  print(train_dog_fnames[:10])

```

 ['cat.590.jpg', 'cat.658.jpg', 'cat.948.jpg', 'cat.226.jpg', 'cat.573.jpg', 'cat.171.jpg', 'cat.165.jpg', 'cat.514.jpg', 'cat.260.jpg', 'cat.937.jpg', 'dog.828.jpg', 'dog.168.jpg', 'dog.609.jpg', 'dog.446.jpg', 'dog.892.jpg', 'dog.905.jpg', 'dog.521.jpg', 'dog.319.jpg', 'dog.277.jpg', 'dog.771.jpg']

Let's find out the total number of cat and dog images in the train and validation directories:

```

1  print('total training cat images:', len(os.listdir( train_cats_dir )))
2  print('total training dog images:', len(os.listdir( train_dogs_dir )))
3

```

```
4 print('total validation cat images :', len(os.listdir( validation_cats_dir ) )
5 print('total validation dog images :', len(os.listdir( validation_dogs_dir ) )

total training cat images : 1000
total training dog images : 1000
total validation cat images : 500
total validation dog images : 500
```

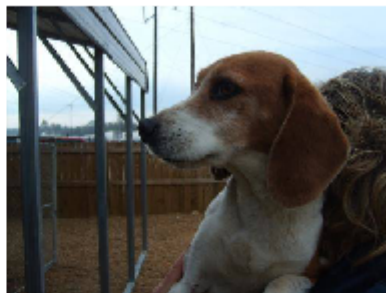
For both cats and dogs, we have 1,000 training images and 500 validation images.

Now let's take a look at a few pictures to get a better sense of what the cat and dog datasets look

```
1 %matplotlib inline
2
3 import matplotlib.image as mpimg
4 import matplotlib.pyplot as plt
5
6 # Parameters for our graph; we'll output images in a 4x4 configuration
7 nrows = 4
8 ncols = 4
9
10 pic_index = 0 # Index for iterating over images
```

Now, display a batch of 8 cat and 8 dog pictures. You can rerun the cell to see a fresh batch each

```
1 # Set up matplotlib fig, and size it to fit 4x4 pics
2 fig = plt.gcf()
3 fig.set_size_inches(ncols*4, nrows*4)
4
5 pic_index+=8
6
7 next_cat_pix = [os.path.join(train_cats_dir, fname)
8                 for fname in train_cat_fnames[ pic_index-8:pic_index]
9                 ]
10
11 next_dog_pix = [os.path.join(train_dogs_dir, fname)
12                 for fname in train_dog_fnames[ pic_index-8:pic_index]
13                 ]
14
15 for i, img_path in enumerate(next_cat_pix+next_dog_pix):
16     # Set up subplot; subplot indices start at 1
17     sp = plt.subplot(nrows, ncols, i + 1)
18     sp.axis('Off') # Don't show axes (or gridlines)
19
20     img = mpimg.imread(img_path)
21     plt.imshow(img)
22
23 plt.show()
24
```



It may not be obvious from looking at the images in this grid, but an important note here, and a sign is that these images come in all shapes and sizes. When you did the handwriting recognition experiment with. These are color and in a variety of shapes. Before training a Neural network with them you'll see the next section.

Ok, now that you have an idea for what your data looks like, the next step is to define the model that will process these images

## ▼ Building a Small Model from Scratch to Get to ~72% Accuracy

In the previous section you saw that the images were in a variety of shapes and sizes. In order to process them, we need them to be in a uniform size. We've chosen 150x150 for this, and you'll see the code that preprocesses the images. But before we continue, let's start defining the model:

But before we continue, let's start defining the model:

Step 1 will be to import tensorflow.

```
1 import tensorflow as tf
```



The default version of TensorFlow in Colab will soon switch to TensorFlow 2.x. We recommend you [upgrade](#) now or ensure your notebook will continue to use TensorFlow 1.x via the %tensorflow\_version --1.15.0 magic command.

Next we will define a Sequential layer as before, adding some convolutional layers first. Note the input shape in the example it was 28x28x1, because the image was 28x28 in greyscale (8 bits, 1 byte for color depth (24 bits, 3 bytes) for the color depth).

We then add a couple of convolutional layers as in the previous example, and flatten the final result.

Finally we add the densely connected layers.

Note that because we are facing a two-class classification problem, i.e. a *binary classification problem*, we need a [sigmoid activation](#), so that the output of our network will be a single scalar between 0 and 1, encoding the probability of class 1 (as opposed to class 0).

```
1 model = tf.keras.models.Sequential([
2     # Note the input shape is the desired size of the image 150x150 with 3 by
3     tf.keras.layers.Conv2D(16, (3,3), activation='relu', input_shape=(150, 150, 3)),
4     tf.keras.layers.MaxPooling2D(2,2),
5     tf.keras.layers.Conv2D(32, (3,3), activation='relu'),
6     tf.keras.layers.MaxPooling2D(2,2),
7     tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
8     tf.keras.layers.MaxPooling2D(2,2),
9     # Flatten the results to feed into a DNN
10    tf.keras.layers.Flatten(),
11    # 512 neuron hidden layer
12    tf.keras.layers.Dense(512, activation='relu'),
13    # Only 1 output neuron. It will contain a value from 0-1 where 0 for 1 class
14    tf.keras.layers.Dense(1, activation='sigmoid')
15 ])
```



WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow\_core/ops/nn/conv\_ops.py:113: Instructions for updating:  
If using Keras pass \*\_constraint arguments to layers.

The `model.summary()` method call prints a summary of the NN

```
1 model.summary()
```



Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 148, 148, 16)	448
max_pooling2d (MaxPooling2D)	(None, 74, 74, 16)	0
conv2d_1 (Conv2D)	(None, 72, 72, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 36, 36, 32)	0
conv2d_2 (Conv2D)	(None, 34, 34, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 17, 17, 64)	0
flatten (Flatten)	(None, 18496)	0
dense (Dense)	(None, 512)	9470464
dense_1 (Dense)	(None, 1)	513
Total params: 9,494,561		
Trainable params: 9,494,561		
Non-trainable params: 0		

The "output shape" column shows how the size of your feature map evolves in each successive layer. The feature maps by a bit due to padding, and each pooling layer halves the dimensions.

Next, we'll configure the specifications for model training. We will train our model with the binary classification problem and our final activation is a sigmoid. (For a refresher on loss metrics, see the [tutorial](#)). We will use the `rmsprop` optimizer with a learning rate of `0.001`. During training, we will want to monitor

**NOTE:** In this case, using the [RMSprop optimization algorithm](#) is preferable to [stochastic gradient descent](#) for learning-rate tuning for us. (Other optimizers, such as [Adam](#) and [Adagrad](#), also automatically adapt their learning rates and work equally well here.)

```
1 from tensorflow.keras.optimizers import RMSprop
2
3 model.compile(optimizer=RMSprop(lr=0.001),
4               loss='binary_crossentropy',
5               metrics = ['acc'])
```



## ▼ Data Preprocessing

Let's set up data generators that will read pictures in our source folders, convert them to float32 our network. We'll have one generator for the training images and one for the validation images. Of size 150x150 and their labels (binary).

As you may already know, data that goes into neural networks should usually be normalized in some processing by the network. (It is uncommon to feed raw pixels into a convnet.) In our case, we will pixel values to be in the  $[0, 1]$  range (originally all values are in the  $[0, 255]$  range).

In Keras this can be done via the `keras.preprocessing.image.ImageDataGenerator` class using the `ImageDataGenerator` class allows you to instantiate generators of augmented image batches (`flow_from_directory(directory)`). These generators can then be used with the Keras model inputs: `fit_generator`, `evaluate_generator`, and `predict_generator`.

```

1  from tensorflow.keras.preprocessing.image import ImageDataGenerator
2
3  # All images will be rescaled by 1./255.
4  train_datagen = ImageDataGenerator( rescale = 1.0/255. )
5  test_datagen  = ImageDataGenerator( rescale = 1.0/255. )
6
7  # -----
8  # Flow training images in batches of 20 using train_datagen generator
9  # -----
10 train_generator = train_datagen.flow_from_directory(train_dir,
11                                                    batch_size=20,
12                                                    class_mode='binary',
13                                                    target_size=(150, 150))
14 # -----
15 # Flow validation images in batches of 20 using test_datagen generator
16 # -----
17 validation_generator = test_datagen.flow_from_directory(validation_dir,
18                                                         batch_size=20,
19                                                         class_mode = 'binary',
20                                                         target_size = (150,
21
```



Found 2000 images belonging to 2 classes.  
Found 1000 images belonging to 2 classes.

## ▼ Training

Let's train on all 2,000 images available, for 15 epochs, and validate on all 1,000 test images. (This will take a while). Do note the values per epoch.

You'll see 4 values per epoch -- Loss, Accuracy, Validation Loss and Validation Accuracy.

The Loss and Accuracy are a great indication of progress of training. It's making a guess as to the correct label by measuring it against the known label, calculating the result. Accuracy is the portion of correct guesses.

measurement with the data that has not been used in training. As expected this would be a bit low section on overfitting later in this course.

```
1 history = model.fit_generator(train_generator,
2                               validation_data=validation_generator,
3                               steps_per_epoch=100,
4                               epochs=15,
5                               validation_steps=50,
6                               verbose=2)
```



```
Epoch 1/15
Epoch 1/15
100/100 - 15s - loss: 0.8980 - acc: 0.5570 - val_loss: 0.6601 - val_acc: 0.60
Epoch 2/15
Epoch 1/15
100/100 - 8s - loss: 0.6553 - acc: 0.6590 - val_loss: 0.6102 - val_acc: 0.680
Epoch 3/15
Epoch 1/15
100/100 - 8s - loss: 0.5711 - acc: 0.7035 - val_loss: 0.5886 - val_acc: 0.707
Epoch 4/15
Epoch 1/15
100/100 - 8s - loss: 0.4989 - acc: 0.7645 - val_loss: 0.5833 - val_acc: 0.719
Epoch 5/15
Epoch 1/15
100/100 - 8s - loss: 0.4263 - acc: 0.7965 - val_loss: 0.6079 - val_acc: 0.712
Epoch 6/15
Epoch 1/15
100/100 - 8s - loss: 0.3507 - acc: 0.8520 - val_loss: 0.6109 - val_acc: 0.717
Epoch 7/15
Epoch 1/15
100/100 - 8s - loss: 0.2767 - acc: 0.8875 - val_loss: 0.7339 - val_acc: 0.717
Epoch 8/15
Epoch 1/15
100/100 - 8s - loss: 0.1884 - acc: 0.9240 - val_loss: 0.8298 - val_acc: 0.713
Epoch 9/15
Epoch 1/15
100/100 - 8s - loss: 0.1268 - acc: 0.9540 - val_loss: 1.1452 - val_acc: 0.718
Epoch 10/15
Epoch 1/15
100/100 - 8s - loss: 0.1003 - acc: 0.9735 - val_loss: 1.0673 - val_acc: 0.719
Epoch 11/15
Epoch 1/15
100/100 - 8s - loss: 0.0640 - acc: 0.9815 - val_loss: 1.3476 - val_acc: 0.713
Epoch 12/15
Epoch 1/15
100/100 - 8s - loss: 0.1177 - acc: 0.9820 - val_loss: 1.4038 - val_acc: 0.706
Epoch 13/15
Epoch 1/15
100/100 - 8s - loss: 0.0407 - acc: 0.9845 - val_loss: 1.9012 - val_acc: 0.713
Epoch 14/15
Epoch 1/15
100/100 - 8s - loss: 0.0803 - acc: 0.9845 - val_loss: 1.7132 - val_acc: 0.707
Epoch 15/15
Epoch 1/15
100/100 - 8s - loss: 0.0505 - acc: 0.9865 - val_loss: 2.2358 - val_acc: 0.703
```



## ▼ Running the Model

Let's now take a look at actually running a prediction using the model. This code will allow you to upload images, which will then be processed by the model, giving an indication of whether the object is a cat or a dog.

```

1  import numpy as np
2
3  from google.colab import files
4  from keras.preprocessing import image
5
6  uploaded=files.upload()
7
8  for fn in uploaded.keys():
9
10     # predicting images
11     path='/content/' + fn
12     img=image.load_img(path, target_size=(150, 150))
13
14     x=image.img_to_array(img)
15     x=np.expand_dims(x, axis=0)
16     images = np.vstack([x])
17
18     classes = model.predict(images, batch_size=10)
19
20     print(classes[0])
21
22     if classes[0]>0:
23         print(fn + " is a dog")
24
25     else:
26         print(fn + " is a cat")
27

```



Choose Files download.jpeg

• **download.jpeg**(image/jpeg) - 6581 bytes, last modified: 16/03/2020 - 100% done  
 Saving download.jpeg to download.jpeg  
 [0.]  
 download.jpeg is a cat

## ▼ Visualizing Intermediate Representations

To get a feel for what kind of features our convnet has learned, one fun thing to do is to visualize the output of the convnet through the convnet.

Let's pick a random cat or dog image from the training set, and then generate a figure where each row is a specific filter in that output feature map. Rerun this cell to generate intermediate representations.

```

1  import numpy as np
2  import random
3  from tensorflow.keras.preprocessing.image import img_to_array, load_img
4
5  # Let's define a new Model that will take an image as input, and will output
6  # intermediate representations for all layers in the previous model after
7  # the first

```

```

/ # the list.
8 successive_outputs = [layer.output for layer in model.layers[1:]]
9
10 #visualization_model = Model(img_input, successive_outputs)
11 visualization_model = tf.keras.models.Model(inputs = model.input, outputs = s
12
13 # Let's prepare a random input image of a cat or dog from the training set.
14 cat_img_files = [os.path.join(train_cats_dir, f) for f in train_cat_fnames]
15 dog_img_files = [os.path.join(train_dogs_dir, f) for f in train_dog_fnames]
16
17 img_path = random.choice(cat_img_files + dog_img_files)
18 img = load_img(img_path, target_size=(150, 150)) # this is a PIL image
19
20 x = img_to_array(img) # Numpy array with shape (1
21 x = x.reshape((1,) + x.shape) # Numpy array with shape (1
22
23 # Rescale by 1/255
24 x /= 255.0
25
26 # Let's run our image through our network, thus obtaining all
27 # intermediate representations for this image.
28 successive_feature_maps = visualization_model.predict(x)
29
30 # These are the names of the layers, so can have them as part of our plot
31 layer_names = [layer.name for layer in model.layers]
32
33 # -----
34 # Now let's display our representations
35 # -----
36 for layer_name, feature_map in zip(layer_names, successive_feature_maps):
37
38     if len(feature_map.shape) == 4:
39
40         #-----
41         # Just do this for the conv / maxpool layers, not the fully-connected lay
42         #-----
43         n_features = feature_map.shape[-1] # number of features in the feature m
44         size = feature_map.shape[1] # feature map shape (1, size, size, n
45
46         # We will tile our images in this matrix
47         display_grid = np.zeros((size, size * n_features))
48
49         #-----
50         # Postprocess the feature to be visually palatable
51         #-----
52         for i in range(n_features):
53             x = feature_map[0, :, :, i]
54             x -= x.mean()
55             x /= x.std ()
56             x *= 64
57             x += 128
58             x = np.clip(x, 0, 255).astype('uint8')
59             display_grid[:, i * size : (i + 1) * size] = x # Tile each filter into
60
61         #-----
62         # Display the grid

```

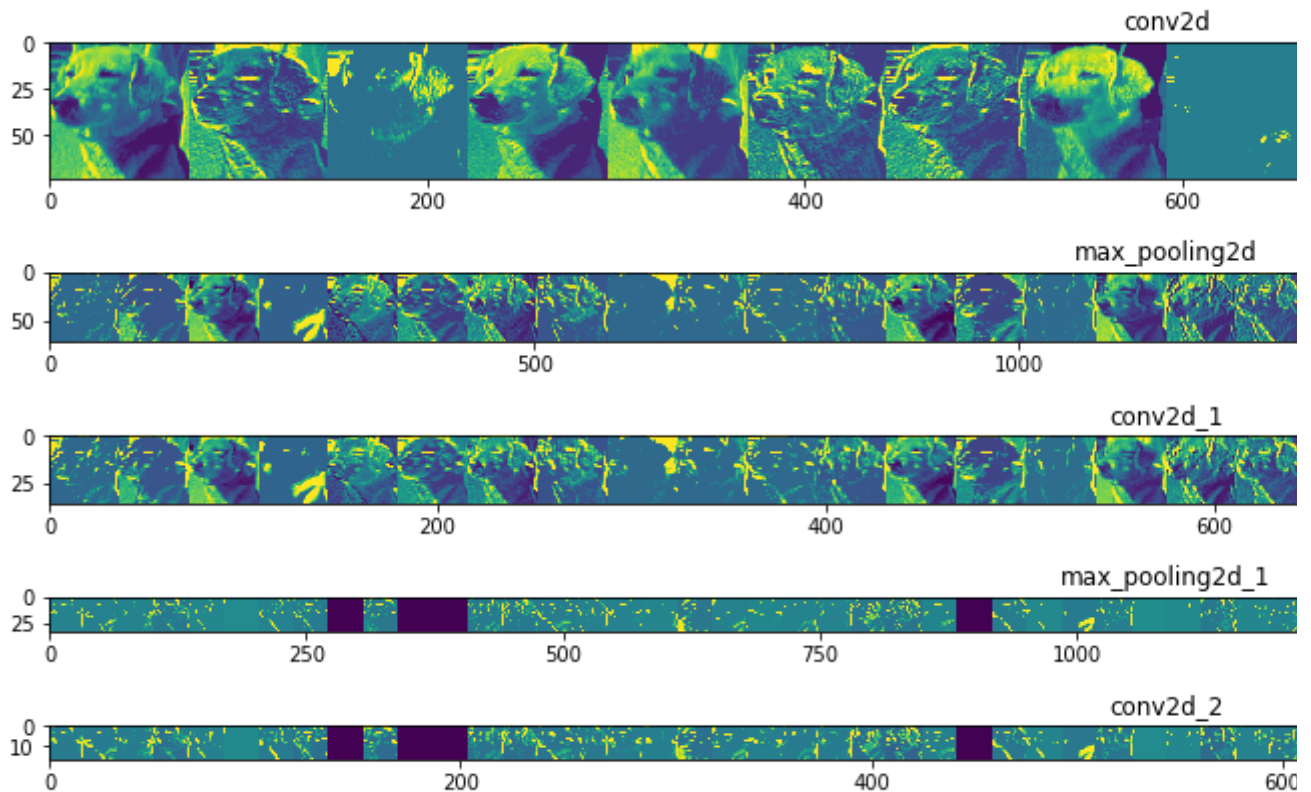
```

62     # Display the grid
63     #-----
64
65     scale = 20. / n_features
66     plt.figure( figsize=(scale * n_features, scale) )
67     plt.title ( layer_name )
68     plt.grid ( False )
69     plt.imshow( display_grid, aspect='auto', cmap='viridis' )

```



/usr/local/lib/python3.6/dist-packages/ipykernel\_launcher.py:55: RuntimeWarni



As you can see we go from the raw pixels of the images to increasingly abstract and compact representations. Downstream starts highlighting what the network pays attention to, and they show fewer and fewer zero. This is called "sparsity." Representation sparsity is a key feature of deep learning.

These representations carry increasingly less information about the original pixels of the image, but more information about the class of the image. You can think of a convnet (or a deep network in general) as an information distiller.

## ▼ Evaluating Accuracy and Loss for the Model

Let's plot the training/validation accuracy and loss as collected during training:

```

1  #-----
2  # Retrieve a list of list results on training and test data
3  # sets for each training epoch
4  #-----
5  acc      = history.history[ 'acc' ]
6  val_acc  = history.history[ 'val_acc' ]
7  loss     = history.history[ 'loss' ]
8  val_loss = history.history[ 'val_loss' ]
9
10 epochs   = range(len(acc)) # Get number of epochs

```

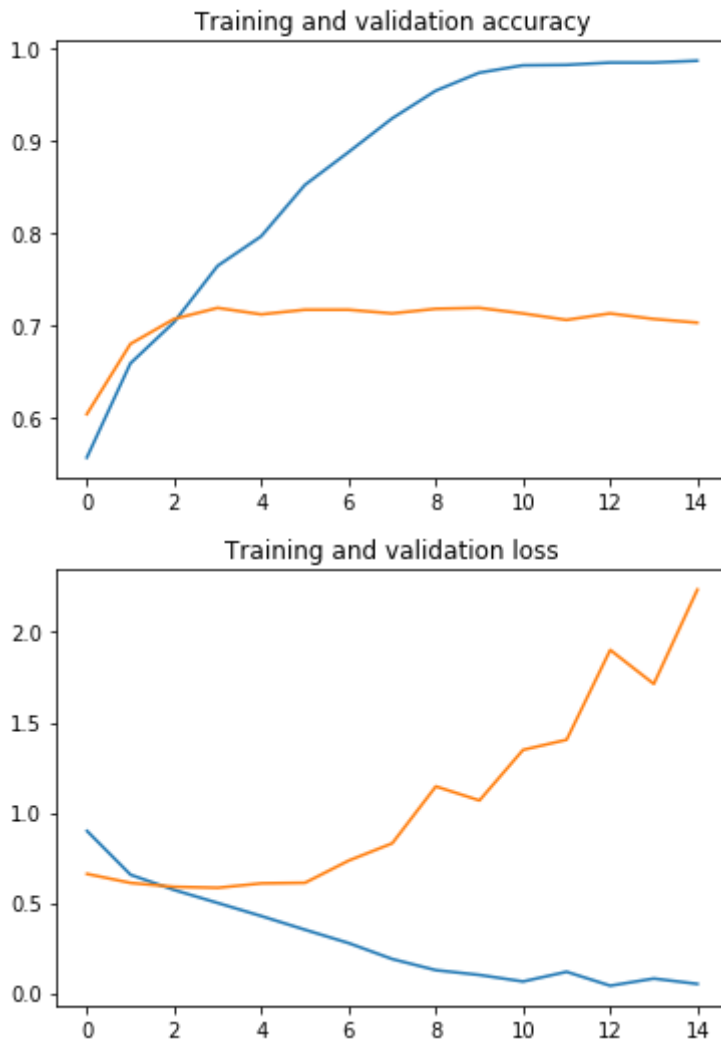
```

10 epochs = range(len(acc)) # get number of epochs
11
12 #-----
13 # Plot training and validation accuracy per epoch
14 #-----
15 plt.plot ( epochs, acc )
16 plt.plot ( epochs, val_acc )
17 plt.title ('Training and validation accuracy')
18 plt.figure()
19
20 #-----
21 # Plot training and validation loss per epoch
22 #-----
23 plt.plot ( epochs, loss )
24 plt.plot ( epochs, val_loss )
25 plt.title ('Training and validation loss' )

```

--VISUAL--

 Text(0.5, 1.0, 'Training and validation loss')



As you can see, we are **overfitting** like it's getting out of fashion. Our training accuracy (in blue) goes up, but our validation accuracy (in green) stalls at 70%. Our validation loss reaches its minimum after only five epochs.

Since we have a relatively small number of training examples (2000), overfitting should be our main concern. Our model, exposed to too few examples, learns patterns that do not generalize to new data, i.e. when making predictions. For instance, if you, as a human, only see three images of people who are running, you might learn that all people who are running are wearing a red shirt.

sailors, and among them the only person wearing a cap is a lumberjack, you might start thinking the lumberjack as opposed to a sailor. You would then make a pretty lousy lumberjack/sailor classifier.

Overfitting is the central problem in machine learning: given that we are fitting the parameters of a model, how can we make sure that the representations learned by the model will be applicable to data never seen before? Is a model specific to the training data?

In the next exercise, we'll look at ways to prevent overfitting in the cat vs. dog classification model.

## ▼ Clean Up

Before running the next exercise, run the following cell to terminate the kernel and free memory resources.

```
1 import os, signal
2
3 os.kill(os.getpid(),
4         signal.SIGKILL)
5
```