

```

1 !wget --no-check-certificate \
2   https://storage.googleapis.com/laurencemoroney-blog.appspot.com/horse-or-
3   -0 /tmp/horse-or-human.zip

```

 --2020-03-12 18:04:22-- <https://storage.googleapis.com/laurencemoroney-blog.appspot.com/horse-or-human.zip>
 Resolving storage.googleapis.com (storage.googleapis.com)... 172.217.14.112, 2
 Connecting to storage.googleapis.com (storage.googleapis.com)|172.217.14.112|:
 HTTP request sent, awaiting response... 200 OK
 Length: 149574867 (143M) [application/zip]
 Saving to: '/tmp/horse-or-human.zip'

/tmp/horse-or-human 100%[=====>] 142.65M 101MB/s in 1.4s

2020-03-12 18:04:24 (101 MB/s) - '/tmp/horse-or-human.zip' saved [149574867/14

```

1 !wget --no-check-certificate \
2   https://storage.googleapis.com/laurencemoroney-blog.appspot.com/validation
3   -0 /tmp/validation-horse-or-human.zip

```

 --2020-03-12 18:04:30-- <https://storage.googleapis.com/laurencemoroney-blog.appspot.com/validation-horse-or-human.zip>
 Resolving storage.googleapis.com (storage.googleapis.com)... 172.217.11.176, 2
 Connecting to storage.googleapis.com (storage.googleapis.com)|172.217.11.176|:
 HTTP request sent, awaiting response... 200 OK
 Length: 11480187 (11M) [application/zip]
 Saving to: '/tmp/validation-horse-or-human.zip'

/tmp/validation-hor 100%[=====>] 10.95M --.-KB/s in 0.03s

2020-03-12 18:04:30 (376 MB/s) - '/tmp/validation-horse-or-human.zip' saved [1

The following python code will use the OS library to use Operating System libraries, giving you access to the files, allowing you to unzip the data.

```

1 import os
2 import zipfile
3
4 local_zip = '/tmp/horse-or-human.zip'
5 zip_ref = zipfile.ZipFile(local_zip, 'r')
6 zip_ref.extractall('/tmp/horse-or-human')
7 local_zip = '/tmp/validation-horse-or-human.zip'
8 zip_ref = zipfile.ZipFile(local_zip, 'r')
9 zip_ref.extractall('/tmp/validation-horse-or-human')
10 zip_ref.close()

```

The contents of the .zip are extracted to the base directory /tmp/horse-or-human, which in turn contains subdirectories.

In short: The training set is the data that is used to tell the neural network model that 'this is what a horse looks like' etc.

One thing to pay attention to in this sample: We do not explicitly label the images as horses or humans. In the example earlier, we had labelled 'this is a 1', 'this is a 7' etc. Later you'll see something called an `ImageGenerator` to read images from subdirectories, and automatically label them from the name of that subdirectory. We have a directory containing a 'horses' directory and a 'humans' one. `ImageGenerator` will label the images accordingly.

Let's define each of these directories:

```

1  # Directory with our training horse pictures
2  train_horse_dir = os.path.join('/tmp/horse-or-human/horses')
3
4  # Directory with our training human pictures
5  train_human_dir = os.path.join('/tmp/horse-or-human/humans')
6
7  # Directory with our validation horse pictures
8  validation_horse_dir = os.path.join('/tmp/validation-horse-or-human/validation')
9
10 # Directory with our validation human pictures
11 validation_human_dir = os.path.join('/tmp/validation-horse-or-human/validation')

```

▼ Building a Small Model from Scratch

But before we continue, let's start defining the model:

Step 1 will be to import tensorflow.

```
1 import tensorflow as tf
```



The default version of TensorFlow in Colab will soon switch to TensorFlow 2.x. We recommend you [upgrade](#) now or ensure your notebook will continue to use TensorFlow 1.x via the `%tensorflow1` magic.

We then add convolutional layers as in the previous example, and flatten the final result to feed into a fully connected layer.

Finally we add the densely connected layers.

Note that because we are facing a two-class classification problem, i.e. a *binary classification problem*, we use the [sigmoid activation](#), so that the output of our network will be a single scalar between 0 and 1, encoding the probability of class 1 (as opposed to class 0).

```

1  model = tf.keras.models.Sequential([
2      # Note the input shape is the desired size of the image 150x150 with 3 bytes
3      # This is the first convolution
4      tf.keras.layers.Conv2D(16, (3,3), activation='relu', input_shape=(150, 150, 3)),
5      tf.keras.layers.MaxPooling2D(2, 2),
6      # The second convolution
7      tf.keras.layers.Conv2D(32, (3,3), activation='relu'),
8      tf.keras.layers.MaxPooling2D(2,2),
9      # The third convolution
10     tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
11     tf.keras.layers.MaxPooling2D(2, 2),

```

```

11     tf.keras.layers.MaxPooling2D(2,2),
12     # The fourth convolution
13     #tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
14     #tf.keras.layers.MaxPooling2D(2,2),
15     # The fifth convolution
16     #tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
17     #tf.keras.layers.MaxPooling2D(2,2),
18     # Flatten the results to feed into a DNN
19     tf.keras.layers.Flatten(),
20     # 512 neuron hidden layer
21     tf.keras.layers.Dense(512, activation='relu'),
22     # Only 1 output neuron. It will contain a value from 0-1 where 0 for 1 class
23     tf.keras.layers.Dense(1, activation='sigmoid')
24 ])
```



WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core/...
 Instructions for updating:
 If using Keras pass *_constraint arguments to layers.

The `model.summary()` method call prints a summary of the NN

```
1 model.summary()
```



Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 148, 148, 16)	448
max_pooling2d (MaxPooling2D)	(None, 74, 74, 16)	0
conv2d_1 (Conv2D)	(None, 72, 72, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 36, 36, 32)	0
conv2d_2 (Conv2D)	(None, 34, 34, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 17, 17, 64)	0
flatten (Flatten)	(None, 18496)	0
dense (Dense)	(None, 512)	9470464
dense_1 (Dense)	(None, 1)	513
Total params: 9,494,561		
Trainable params: 9,494,561		
Non-trainable params: 0		

The "output shape" column shows how the size of your feature map evolves in each successive layer. The feature maps by a bit due to padding, and each pooling layer halves the dimensions.

Next, we'll configure the specifications for model training. We will train our model with the binary classification problem and our final activation is a sigmoid. (For a refresher on loss metrics, see the [use the rmsprop optimizer](#) with a learning rate of 0.001. During training, we will want to monitor

NOTE: In this case, using the [RMSprop optimization algorithm](#) is preferable to [stochastic gradient learning-rate tuning](#) for us. (Other optimizers, such as [Adam](#) and [Adagrad](#), also automatically adap

```
1 from tensorflow.keras.optimizers import RMSprop
2
3 model.compile(loss='binary_crossentropy',
4               optimizer=RMSprop(lr=0.001),
5               metrics=['acc'])
```



WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow_core/Instructions for updating:
Use tf.where in 2.0, which has the same broadcast rule as np.where

▼ Data Preprocessing

Let's set up data generators that will read pictures in our source folders, convert them to float32 for our network. We'll have one generator for the training images and one for the validation images. Our images are size 300x300 and their labels (binary).

As you may already know, data that goes into neural networks should usually be normalized in some way before being processed by the network. (It is uncommon to feed raw pixels into a convnet.) In our case, we will normalize pixel values to be in the [0, 1] range (originally all values are in the [0, 255] range).

In Keras this can be done via the `keras.preprocessing.image.ImageDataGenerator` class. The `ImageDataGenerator` class allows you to instantiate generators of augmented image batches (a `flow_from_directory(directory)`). These generators can then be used with the Keras model `fit_generator`, `evaluate_generator`, and `predict_generator`.

```
1 from tensorflow.keras.preprocessing.image import ImageDataGenerator
2
3 # All images will be rescaled by 1./255
4 train_datagen = ImageDataGenerator(rescale=1/255)
5 validation_datagen = ImageDataGenerator(rescale=1/255)
6
7 # Flow training images in batches of 128 using train_datagen generator
8 train_generator = train_datagen.flow_from_directory(
9     '/tmp/horse-or-human/', # This is the source directory for training :
10     target_size=(150, 150), # All images will be resized to 150x150
11     batch_size=128,
12     # Since we use binary_crossentropy loss, we need binary labels
13     class_mode='binary')
14
15 # Flow validation images in batches of 128 using train_datagen generator
16 validation_generator = validation_datagen.flow_from_directory(
17     '/tmp/validation-horse-or-human/', # This is the source directory for
18     target_size=(150, 150), # All images will be resized to 150x150
19     batch_size=32,
```

```
20 # Since we use binary_crossentropy loss, we need binary labels
21 class_mode='binary')
```



Found 1027 images belonging to 2 classes.
Found 256 images belonging to 2 classes.

▼ Training

Let's train for 15 epochs -- this may take a few minutes to run.

Do note the values per epoch.

The Loss and Accuracy are a great indication of progress of training. It's making a guess as to the measuring it against the known label, calculating the result. Accuracy is the portion of correct gue

```
1 history = model.fit_generator(
2     train_generator,
3     steps_per_epoch=8,
4     epochs=15,
5     verbose=1,
6     validation_data = validation_generator,
7     validation_steps=8)
```



```
8/8 [=====] - 1s 91ms/step - loss: 0.4881 - acc: 0.82
8/8 [=====] - 9s 1s/step - loss: 2.6206 - acc: 0.5473
Epoch 2/15
6/8 [=====>.....] - ETA: 1s - loss: 0.4942 - acc: 0.7344Epc
8/8 [=====] - 1s 133ms/step - loss: 0.6018 - acc: 0.7
8/8 [=====] - 5s 674ms/step - loss: 0.4540 - acc: 0.7
Epoch 3/15
7/8 [=====>....] - ETA: 0s - loss: 0.6742 - acc: 0.8392Epc
8/8 [=====] - 1s 129ms/step - loss: 1.5070 - acc: 0.6
8/8 [=====] - 5s 657ms/step - loss: 0.6106 - acc: 0.8
Epoch 4/15
7/8 [=====>....] - ETA: 0s - loss: 0.1917 - acc: 0.9252Epc
8/8 [=====] - 1s 136ms/step - loss: 1.3144 - acc: 0.7
8/8 [=====] - 6s 739ms/step - loss: 0.1875 - acc: 0.9
Epoch 5/15
7/8 [=====>....] - ETA: 0s - loss: 0.1785 - acc: 0.9053Epc
8/8 [=====] - 1s 134ms/step - loss: 1.0523 - acc: 0.8
8/8 [=====] - 5s 651ms/step - loss: 0.1678 - acc: 0.9
Epoch 6/15
7/8 [=====>....] - ETA: 0s - loss: 0.0798 - acc: 0.9650Epc
8/8 [=====] - 1s 130ms/step - loss: 1.0680 - acc: 0.8
8/8 [=====] - 5s 653ms/step - loss: 0.0833 - acc: 0.9
Epoch 7/15
7/8 [=====>....] - ETA: 0s - loss: 0.0744 - acc: 0.9650Epc
8/8 [=====] - 1s 133ms/step - loss: 2.5921 - acc: 0.6
8/8 [=====] - 5s 655ms/step - loss: 0.1387 - acc: 0.9
Epoch 8/15
7/8 [=====>....] - ETA: 0s - loss: 0.2048 - acc: 0.8872Epc
8/8 [=====] - 1s 127ms/step - loss: 1.4794 - acc: 0.8
8/8 [=====] - 5s 565ms/step - loss: 0.1853 - acc: 0.8
Epoch 9/15
7/8 [=====>....] - ETA: 0s - loss: 0.0703 - acc: 0.9710Epc
8/8 [=====] - 1s 131ms/step - loss: 1.4551 - acc: 0.7
8/8 [=====] - 6s 729ms/step - loss: 0.0681 - acc: 0.9
Epoch 10/15
7/8 [=====>....] - ETA: 0s - loss: 0.0136 - acc: 1.0000Epc
8/8 [=====] - 1s 127ms/step - loss: 1.9596 - acc: 0.7
8/8 [=====] - 5s 661ms/step - loss: 0.0131 - acc: 1.0
Epoch 11/15
7/8 [=====>....] - ETA: 0s - loss: 0.0135 - acc: 0.9974Epc
8/8 [=====] - 1s 132ms/step - loss: 1.8790 - acc: 0.8
8/8 [=====] - 5s 672ms/step - loss: 0.0123 - acc: 0.9
Epoch 12/15
7/8 [=====>....] - ETA: 0s - loss: 0.0042 - acc: 1.0000Epc
8/8 [=====] - 1s 134ms/step - loss: 2.1995 - acc: 0.8
8/8 [=====] - 6s 728ms/step - loss: 0.0040 - acc: 1.0
Epoch 13/15
7/8 [=====>....] - ETA: 0s - loss: 9.5277e-04 - acc: 1.000
8/8 [=====] - 1s 134ms/step - loss: 2.4533 - acc: 0.8
8/8 [=====] - 5s 575ms/step - loss: 9.6421e-04 - acc:
Epoch 14/15
7/8 [=====>....] - ETA: 0s - loss: 0.3477 - acc: 0.8772Epc
8/8 [=====] - 1s 132ms/step - loss: 1.0312 - acc: 0.7
8/8 [=====] - 6s 732ms/step - loss: 0.3178 - acc: 0.8
Epoch 15/15
7/8 [=====>....] - ETA: 0s - loss: 0.1619 - acc: 0.9416Epc
8/8 [=====] - 1s 131ms/step - loss: 1.7745 - acc: 0.8
8/8 [=====] - 5s 658ms/step - loss: 0.1443 - acc: 0.9
```

▼ Running the Model

Let's now take a look at actually running a prediction using the model. This code will allow you to upload files, which will then upload them, and run them through the model, giving an indication of whether the object

```

1  import numpy as np
2  from google.colab import files
3  from keras.preprocessing import image
4
5  uploaded = files.upload()
6
7  for fn in uploaded.keys():
8
9      # predicting images
10     path = '/content/' + fn
11     img = image.load_img(path, target_size=(150, 150))
12     x = image.img_to_array(img)
13     x = np.expand_dims(x, axis=0)
14
15     images = np.vstack([x])
16     classes = model.predict(images, batch_size=10)
17     print(classes[0])
18     if classes[0]>0.5:
19         print(fn + " is a human")
20     else:
21         print(fn + " is a horse")
22

```



Choose Files jpggggg.jpg

- **jpggggg.jpg**(image/jpeg) - 28499 bytes, last modified: 28/11/2019 - 100% done

Saving jpggggg.jpg to jpggggg.jpg
 [0.]
 jpggggg.jpg is a horse

▼ Visualizing Intermediate Representations

To get a feel for what kind of features our convnet has learned, one fun thing to do is to visualize h through the convnet.

Let's pick a random image from the training set, and then generate a figure where each row is the c a specific filter in that output feature map. Rerun this cell to generate intermediate representations

```

1  import numpy as np
2  import random
3  from tensorflow.keras.preprocessing.image import img_to_array, load_img
4
5  # Let's define a new Model that will take an image as input, and will output
6  # intermediate representations for all layers in the previous model after
7  # the first.
8  successive_outputs = [layer.output for layer in model.layers[1:]]

```

```

9  #visualization_model = Model(img_input, successive_outputs)
10 visualization_model = tf.keras.models.Model(inputs = model.input, outputs = s
11 # Let's prepare a random input image from the training set.
12 horse_img_files = [os.path.join(train_horse_dir, f) for f in train_horse_name
13 human_img_files = [os.path.join(train_human_dir, f) for f in train_human_name
14 img_path = random.choice(horse_img_files + human_img_files)
15
16 img = load_img(img_path, target_size=(300, 300)) # this is a PIL image
17 x = img_to_array(img) # Numpy array with shape (150, 150, 3)
18 x = x.reshape((1,) + x.shape) # Numpy array with shape (1, 150, 150, 3)
19
20 # Rescale by 1/255
21 x /= 255
22
23 # Let's run our image through our network, thus obtaining all
24 # intermediate representations for this image.
25 successive_feature_maps = visualization_model.predict(x)
26
27 # These are the names of the layers, so can have them as part of our plot
28 layer_names = [layer.name for layer in model.layers]
29
30 # Now let's display our representations
31 for layer_name, feature_map in zip(layer_names, successive_feature_maps):
32     if len(feature_map.shape) == 4:
33         # Just do this for the conv / maxpool layers, not the fully-connected laye
34         n_features = feature_map.shape[-1] # number of features in feature map
35         # The feature map has shape (1, size, size, n_features)
36         size = feature_map.shape[1]
37         # We will tile our images in this matrix
38         display_grid = np.zeros((size, size * n_features))
39         for i in range(n_features):
40             # Postprocess the feature to make it visually palatable
41             x = feature_map[0, :, :, i]
42             x -= x.mean()
43             x /= x.std()
44             x *= 64
45             x += 128
46             x = np.clip(x, 0, 255).astype('uint8')
47             # We'll tile each filter into this big horizontal grid
48             display_grid[:, i * size : (i + 1) * size] = x
49         # Display the grid
50         scale = 20. / n_features
51         plt.figure(figsize=(scale * n_features, scale))
52         plt.title(layer_name)
53         plt.grid(False)
54         plt.imshow(display_grid, aspect='auto', cmap='viridis')

```

--NORMAL--

As you can see we go from the raw pixels of the images to increasingly abstract and compact representations. The first few layers downstream start highlighting what the network pays attention to, and they show fewer and fewer zero. This is called "sparsity." Representation sparsity is a key feature of deep learning.

These representations carry increasingly less information about the original pixels of the image, but

▼ Clean Up

Before running the next exercise, run the following cell to terminate the kernel and free memory resources.

```
1 import os, signal
2 os.kill(os.getpid(), signal.SIGKILL)
```