

Below is code with a link to a happy or sad dataset which contains 80 images, 40 happy and 40 sad. Create a convolutional neural network that trains to 100% accuracy on these images, which cancels training upon hitting training accuracy of >.999

Hint -- it will work best with 3 convolutional layers.

In [29]:

```
import tensorflow as tf
import os
import zipfile
from os import path, getcwd, chdir

# DO NOT CHANGE THE LINE BELOW. If you are developing in a local
# environment, then grab happy-or-sad.zip from the Coursera Jupyter Notebook
# and place it inside a local folder and edit the path to that location
path = f"{getcwd()}/../tmp2/happy-or-sad.zip"

zip_ref = zipfile.ZipFile(path, 'r')
zip_ref.extractall("/tmp/h-or-s")
zip_ref.close()
```

In [30]:

```

GRADED FUNCTION: train_happy_sad_model
f train_happy_sad_model():
    # Please write your code only where you are indicated.
    # please do not remove # model fitting inline comments.

    DESIRED_ACCURACY = 0.999

    # YOUR CODE STARTS HERE
    class myCallback(tf.keras.callbacks.Callback):
        def on_epoch_end(self, epoch, logs={}):
            if(logs.get('acc')>0.999):
                print("\nReached 99.9% accuracy so cancelling training!")
                self.model.stop_training = True
    # YOUR CODE ENDS HERE

    callbacks = myCallback()

    # This Code Block should Define and Compile the Model. Please assume the images are
    model = tf.keras.models.Sequential([
    # Note the input shape is the desired size of the image 300x300 with 3 bytes color
    # This is the first convolution
        tf.keras.layers.Conv2D(16, (3,3), activation='relu', input_shape=(300, 300, 3)),
        tf.keras.layers.MaxPooling2D(2, 2),
    # The second convolution
        tf.keras.layers.Conv2D(32, (3,3), activation='relu'),
        tf.keras.layers.MaxPooling2D(2,2),
    # The third convolution
        tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
        tf.keras.layers.MaxPooling2D(2,2),
    # Flatten the results to feed into a DNN
        tf.keras.layers.Flatten(),
    # 512 neuron hidden layer
        tf.keras.layers.Dense(512, activation='relu'),
    # Only 1 output neuron. It will contain a value from 0-1 where 0 for 1 class ('happy')
        tf.keras.layers.Dense(1, activation='sigmoid')
    ])

    from tensorflow.keras.optimizers import RMSprop

    model.compile(loss='binary_crossentropy',
                  optimizer=RMSprop(lr=0.001),
                  metrics=['acc'])

    # This code block should create an instance of an ImageDataGenerator called train_datagen
    # And a train_generator by calling train_datagen.flow_from_directory

    from tensorflow.keras.preprocessing.image import ImageDataGenerator

    train_datagen = ImageDataGenerator(rescale=1/255)# Your Code Here

    # Please use a target_size of 150 X 150.

    train_generator = train_datagen.flow_from_directory(
        "/tmp/h-or-s", # This is the source directory for training images
        target_size=(300, 300), # All images will be resized to 150x150
        batch_size=128,
        # Since we use binary_crossentropy loss, we need binary labels
        class_mode='binary')

```

```
# Your Code Here)
# Expected output: 'Found 80 images belonging to 2 classes'

# This code block should call model.fit_generator and train for
# a number of epochs.
# model fitting
history = model.fit_generator(
    train_generator,
    steps_per_epoch=8,
    epochs=15,
    verbose=1)
    # Your Code Here)
# model fitting
return history.history['acc'][-1]
```

In [31]:

```
# The Expected output: "Reached 99.9% accuracy so cancelling training!"
train_happy_sad_model()
```

```
Found 80 images belonging to 2 classes.
Epoch 1/15
8/8 [=====] - 9s 1s/step - loss: 4.1431 - acc: 0.5312
Epoch 2/15
8/8 [=====] - 5s 675ms/step - loss: 0.3291 - acc: 0.8375
Epoch 3/15
8/8 [=====] - 5s 674ms/step - loss: 0.0786 - acc: 0.9719
Epoch 4/15
8/8 [=====] - 5s 674ms/step - loss: 0.0431 - acc: 0.9891
Epoch 5/15
8/8 [=====] - 5s 662ms/step - loss: 0.0143 - acc: 1.0000
Epoch 6/15
8/8 [=====] - 6s 699ms/step - loss: 0.0084 - acc: 1.0000
Epoch 7/15
8/8 [=====] - 5s 687ms/step - loss: 0.0026 - acc: 1.0000
Epoch 8/15
8/8 [=====] - 5s 687ms/step - loss: 0.0013 - acc: 1.0000
Epoch 9/15
8/8 [=====] - 5s 687ms/step - loss: 7.0013e-04 - acc: 1.0000
Epoch 10/15
8/8 [=====] - 5s 687ms/step - loss: 3.4169e-04 - acc: 1.0000
Epoch 11/15
8/8 [=====] - 5s 687ms/step - loss: 1.7692e-04 - acc: 1.0000
Epoch 12/15
8/8 [=====] - 5s 674ms/step - loss: 9.7270e-05 - acc: 1.0000
Epoch 13/15
8/8 [=====] - 5s 674ms/step - loss: 5.3621e-05 - acc: 1.0000
Epoch 14/15
8/8 [=====] - 6s 699ms/step - loss: 2.9916e-05 - acc: 1.0000
Epoch 15/15
8/8 [=====] - 5s 675ms/step - loss: 1.6872e-05 - acc: 1.0000
```

Out[31]:

1.0

In [4]:

```
# Now click the 'Submit Assignment' button above.
# Once that is complete, please run the following two cells to save your work and c
```

In []:

```
%%javascript
<!-- Save the notebook -->
IPython.notebook.save_checkpoint();
```

In []:

```
%%javascript
<!-- Shutdown and close the notebook -->
window.onbeforeunload = null
window.close();
IPython.notebook.session.delete();
```