

```

1  !wget --no-check-certificate \
2      https://storage.googleapis.com/laurencemoroney-blog.appspot.com/horse-or-
3      -0 /tmp/horse-or-human.zip

--2020-03-12 17:02:29-- https://storage.googleapis.com/laurencemoroney-blog.
Resolving storage.googleapis.com (storage.googleapis.com)... 172.217.214.128,
Connecting to storage.googleapis.com (storage.googleapis.com)|172.217.214.128
HTTP request sent, awaiting response... 200 OK
Length: 149574867 (143M) [application/zip]
Saving to: '/tmp/horse-or-human.zip'

/tmp/horse-or-human 100%[=====>] 142.65M  42.1MB/s   in 3.6s

2020-03-12 17:02:33 (39.8 MB/s) - '/tmp/horse-or-human.zip' saved [149574867/

```

The following python code will use the OS library to use Operating System libraries, giving you access allowing you to unzip the data.

```

1  import os
2  import zipfile
3
4  local_zip = '/tmp/horse-or-human.zip'
5  zip_ref = zipfile.ZipFile(local_zip, 'r')
6  zip_ref.extractall('/tmp/horse-or-human')
7  zip_ref.close()

```

The contents of the .zip are extracted to the base directory /tmp/horse-or-human, which in turn contains subdirectories.

In short: The training set is the data that is used to tell the neural network model that 'this is what I like' etc.

One thing to pay attention to in this sample: We do not explicitly label the images as horses or humans. In the example earlier, we had labelled 'this is a 1', 'this is a 7' etc. Later you'll see something called an ImageGenerator to read images from subdirectories, and automatically label them from the name of that subdirectory. We have a directory containing a 'horses' directory and a 'humans' one. ImageGenerator will label the images accordingly.

Let's define each of these directories:

```

1  # Directory with our training horse pictures
2  train_horse_dir = os.path.join('/tmp/horse-or-human/horses')
3
4  # Directory with our training human pictures
5  train_human_dir = os.path.join('/tmp/horse-or-human/humans')

```

Your session crashed for an unknown reason. [View runtime logs](#) ✕

Now, let's see what the filenames look like in the horses and humans training directories:

```

1  train_horse_names = os.listdir(train_horse_dir)
2  print(train_horse_names[:10])

```

```

3
4 train_human_names = os.listdir(train_human_dir)
5 print(train_human_names[:10])

['horse11-2.png', 'horse27-3.png', 'horse07-1.png', 'horse43-5.png', 'horse01
['human13-14.png', 'human03-20.png', 'human08-26.png', 'human15-28.png', 'hum

```

Let's find out the total number of horse and human images in the directories:

```

1 print('total training horse images:', len(os.listdir(train_horse_dir)))
2 print('total training human images:', len(os.listdir(train_human_dir)))

total training horse images: 500
total training human images: 527

```

Now let's take a look at a few pictures to get a better sense of what they look like. First, configure

```

1 %matplotlib inline
2
3 import matplotlib.pyplot as plt
4 import matplotlib.image as mpimg
5
6 # Parameters for our graph; we'll output images in a 4x4 configuration
7 nrows = 4
8 ncols = 4
9
10 # Index for iterating over images
11 pic_index = 0

```

Now, display a batch of 8 horse and 8 human pictures. You can rerun the cell to see a fresh batch

```

1 # Set up matplotlib fig, and size it to fit 4x4 pics
2 fig = plt.gcf()
3 fig.set_size_inches(ncols * 4, nrows * 4)
4
5 pic_index += 8
6 next_horse_pix = [os.path.join(train_horse_dir, fname)
7                   for fname in train_horse_names[pic_index-8:pic_index]]
8 next_human_pix = [os.path.join(train_human_dir, fname)
9                  for fname in train_human_names[pic_index-8:pic_index]]
10
11 for i, img_path in enumerate(next_horse_pix+next_human_pix):
12     # Set up subplot; subplot indices start at 1
13     sp = plt.subplot(nrows, ncols, i + 1)
14     sp.axis('off') # Don't show axes (or gridlines)
17     plt.imshow(img)
18
19 plt.show()
20

```

Your session crashed for an unknown reason. [View runtime logs](#) ✕



Your session crashed for an unknown reason. [View runtime logs](#) ✕

But before we continue, let's start defining the model:

Step 1 will be to import tensorflow.

```
1 import tensorflow as tf
```



The default version of TensorFlow in Colab will soon switch to TensorFlow 2.x. We recommend you [upgrade](#) now or ensure your notebook will continue to use TensorFlow 1.x via the `%tensorflow_version 1.x` magic: [more info](#).

We then add convolutional layers as in the previous example, and flatten the final result to feed in

Finally we add the densely connected layers.

Note that because we are facing a two-class classification problem, i.e. a *binary classification problem*, so that the output of our network will be a single scalar between 0 and 1, encoding the (as opposed to class 0).

```
1 model = tf.keras.models.Sequential([
2     # Note the input shape is the desired size of the image 300x300 with 3 by
3     # This is the first convolution
4     tf.keras.layers.Conv2D(16, (3,3), activation='relu', input_shape=(300, 300, 3)),
5     tf.keras.layers.MaxPooling2D(2, 2),
6     # The second convolution
7     tf.keras.layers.Conv2D(32, (3,3), activation='relu'),
8     tf.keras.layers.MaxPooling2D(2,2),
9     # The third convolution
10    tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
11    tf.keras.layers.MaxPooling2D(2,2),
12    # The fourth convolution
13    tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
14    tf.keras.layers.MaxPooling2D(2,2),
15    # The fifth convolution
16    tf.keras.layers.Conv2D(64, (3,3), activation='relu'),
17    tf.keras.layers.MaxPooling2D(2,2),
18    # Flatten the results to feed into a DNN
19    tf.keras.layers.Flatten(),
20    # 512 neuron hidden layer
21    tf.keras.layers.Dense(512, activation='relu'),
22    # Only 1 output neuron. It will contain a value from 0-1 where 0 for 1 class and 1 for the other class
23    tf.keras.layers.Dense(1, activation='sigmoid')
24 ])
```



WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow\_core/keras/layers/convolutional.py:106: The name tf.nn.conv2d is deprecated. Please use tf.nn.conv2d\_v2 instead.  
Instructions for updating:  
If using Keras pass \*\_constraint arguments to layers.

The `model.summary()` method call prints a summary of the NN

Your session crashed for an unknown reason. [View runtime logs](#) ✕



Model: "sequential"

Layer (type)	Output Shape	Param #
conv2d (Conv2D)	(None, 298, 298, 16)	448
max_pooling2d (MaxPooling2D)	(None, 149, 149, 16)	0
conv2d_1 (Conv2D)	(None, 147, 147, 32)	4640
max_pooling2d_1 (MaxPooling2D)	(None, 73, 73, 32)	0
conv2d_2 (Conv2D)	(None, 71, 71, 64)	18496
max_pooling2d_2 (MaxPooling2D)	(None, 35, 35, 64)	0
conv2d_3 (Conv2D)	(None, 33, 33, 64)	36928
max_pooling2d_3 (MaxPooling2D)	(None, 16, 16, 64)	0
conv2d_4 (Conv2D)	(None, 14, 14, 64)	36928
max_pooling2d_4 (MaxPooling2D)	(None, 7, 7, 64)	0
flatten (Flatten)	(None, 3136)	0
dense (Dense)	(None, 512)	1606144
dense_1 (Dense)	(None, 1)	513
Total params: 1,704,097		
Trainable params: 1,704,097		
Non-trainable params: 0		

The "output shape" column shows how the size of your feature map evolves in each successive layer. The feature maps by a bit due to padding, and each pooling layer halves the dimensions.

Next, we'll configure the specifications for model training. We will train our model with the binary classification problem and our final activation is a sigmoid. (For a refresher on loss metrics, see the [loss metrics](#) page.) We will use the `rmsprop` optimizer with a learning rate of `0.001`. During training, we will want to monitor

**NOTE:** In this case, using the [RMSprop optimization algorithm](#) is preferable to [stochastic gradient descent](#) for us. (Other optimizers, such as [Adam](#) and [Adagrad](#), also automatically adapt their learning-rate tuning for us.)

```
1 from tensorflow.keras.optimizers import RMSprop
```

Your session crashed for an unknown reason. [View runtime logs](#) ✕

```
5         metrics=['acc'])
```



WARNING:tensorflow:From /usr/local/lib/python3.6/dist-packages/tensorflow\_core/ops/nn/nn\_ops\_inplace\_merge\_grads\_v2.py:11: nn\_ops\_inplace\_merge\_grads\_v2 is deprecated and will be removed in a future version. Instructions for updating:  
Use tf.where in 2.0, which has the same broadcast rule as np.where

## ▼ Data Preprocessing

Let's set up data generators that will read pictures in our source folders, convert them to float32 for our network. We'll have one generator for the training images and one for the validation images. Our images are size 300x300 and their labels (binary).

As you may already know, data that goes into neural networks should usually be normalized in some way before processing by the network. (It is uncommon to feed raw pixels into a convnet.) In our case, we will normalize pixel values to be in the  $[0, 1]$  range (originally all values are in the  $[0, 255]$  range).

In Keras this can be done via the `keras.preprocessing.image.ImageDataGenerator` class using the `ImageDataGenerator` class allows you to instantiate generators of augmented image batches (`flow_from_directory(directory)`). These generators can then be used with the Keras model `fit_generator`, `evaluate_generator`, and `predict_generator`.

```
1 from tensorflow.keras.preprocessing.image import ImageDataGenerator
2
3 # All images will be rescaled by 1./255
4 train_datagen = ImageDataGenerator(rescale=1/255)
5
6 # Flow training images in batches of 128 using train_datagen generator
7 train_generator = train_datagen.flow_from_directory(
8     '/tmp/horse-or-human/', # This is the source directory for training
9     target_size=(300, 300), # All images will be resized to 150x150
10    batch_size=128,
11    # Since we use binary_crossentropy loss, we need binary labels
12    class_mode='binary')
13
```

 Found 1027 images belonging to 2 classes.

## ▼ Training

Let's train for 15 epochs -- this may take a few minutes to run.

Do note the values per epoch.

The Loss and Accuracy are a great indication of progress of training. It's making a guess as to the result by measuring it against the known label, calculating the result. Accuracy is the portion of correct guesses.

Your session crashed for an unknown reason. [View runtime logs](#) ✕

```
1 history = model.fit_generator(
2     train_generator,
3     steps_per_epoch=8,
4     epochs=15,
5     verbose=1)
```



```

Epoch 1/15
8/8 [=====] - 9s 1s/step - loss: 0.6901 - acc: 0.564
Epoch 2/15
8/8 [=====] - 5s 654ms/step - loss: 0.7063 - acc: 0.
Epoch 3/15
8/8 [=====] - 6s 743ms/step - loss: 0.7360 - acc: 0.
Epoch 4/15
8/8 [=====] - 6s 739ms/step - loss: 0.5188 - acc: 0.
Epoch 5/15
8/8 [=====] - 6s 741ms/step - loss: 0.2582 - acc: 0.
Epoch 6/15
8/8 [=====] - 7s 849ms/step - loss: 0.1777 - acc: 0.
Epoch 7/15
8/8 [=====] - 6s 735ms/step - loss: 0.1574 - acc: 0.
Epoch 8/15
8/8 [=====] - 5s 641ms/step - loss: 0.2443 - acc: 0.
Epoch 9/15
8/8 [=====] - 7s 843ms/step - loss: 0.2698 - acc: 0.
Epoch 10/15
8/8 [=====] - 6s 732ms/step - loss: 0.0935 - acc: 0.
Epoch 11/15
8/8 [=====] - 6s 758ms/step - loss: 0.3006 - acc: 0.
Epoch 12/15
8/8 [=====] - 6s 746ms/step - loss: 0.1476 - acc: 0.
Epoch 13/15
8/8 [=====] - 7s 848ms/step - loss: 0.0371 - acc: 0.
Epoch 14/15
8/8 [=====] - 6s 750ms/step - loss: 0.0352 - acc: 0.
Epoch 15/15
8/8 [=====] - 5s 637ms/step - loss: 0.0783 - acc: 0.

```

## ▼ Running the Model

Let's now take a look at actually running a prediction using the model. This code will allow you to upload images, which you can then upload them, and run them through the model, giving an indication of whether the object

```

1  import numpy as np
2  from google.colab import files
3  from keras.preprocessing import image
4
5  uploaded = files.upload()
6
7  for fn in uploaded.keys():
8
9      # predicting images
10     path = '/content/' + fn
11     img = image.load_img(path, target_size=(300, 300))
12     x = image.img_to_array(img)

```

Your session crashed for an unknown reason. [View runtime logs](#) ✕

```

16     classes = model.predict(images, batch_size=10)
17     print(classes[0])
18     if classes[0]>0.5:
19         print(fn + " is a human")

```

```

20     else:
21         print(fn + " is a horse")
22

```



Using TensorFlow backend.

Choose Files finall.jpg

- **finall.jpg**(image/jpeg) - 7249 bytes, last modified: 28/11/2019 - 100% done

Saving finall.jpg to finall.jpg  
[0.9988802]  
finall.jpg is a human

## ▼ Visualizing Intermediate Representations

To get a feel for what kind of features our convnet has learned, one fun thing to do is to visualize it through the convnet.

Let's pick a random image from the training set, and then generate a figure where each row is the a specific filter in that output feature map. Rerun this cell to generate intermediate representation:

```

1  import numpy as np
2  import random
3  from tensorflow.keras.preprocessing.image import img_to_array, load_img
4
5  # Let's define a new Model that will take an image as input, and will output
6  # intermediate representations for all layers in the previous model after
7  # the first.
8  successive_outputs = [layer.output for layer in model.layers[1:]]
9  #visualization_model = Model(img_input, successive_outputs)
10 visualization_model = tf.keras.models.Model(inputs = model.input, outputs = s
11 # Let's prepare a random input image from the training set.
12 horse_img_files = [os.path.join(train_horse_dir, f) for f in train_horse_name
13 human_img_files = [os.path.join(train_human_dir, f) for f in train_human_name
14 img_path = random.choice(horse_img_files + human_img_files)
15
16 img = load_img(img_path, target_size=(300, 300)) # this is a PIL image
17 x = img_to_array(img) # Numpy array with shape (150, 150, 3)
18 x = x.reshape((1,) + x.shape) # Numpy array with shape (1, 150, 150, 3)
19
20 # Rescale by 1/255
21 x /= 255
22
23 # Let's run our image through our network, thus obtaining all
24 # intermediate representations for this image.
25 successive_feature_maps = visualization_model.predict(x)
26
27 # These are the names of the layers, so can have them as part of our plot
28 layer_names = [layer.name for layer in model.layers]

```

Your session crashed for an unknown reason. [View runtime logs](#) ✕

```

31 for layer_name, feature_map in zip(layer_names, successive_feature_maps):
32     if len(feature_map.shape) == 4:
33         # Just do this for the conv / maxpool layers, not the fully-connected lay
34         n_features = feature_map.shape[-1] # number of features in feature map
35         # The feature map has shape (1, size, size, n features)

```

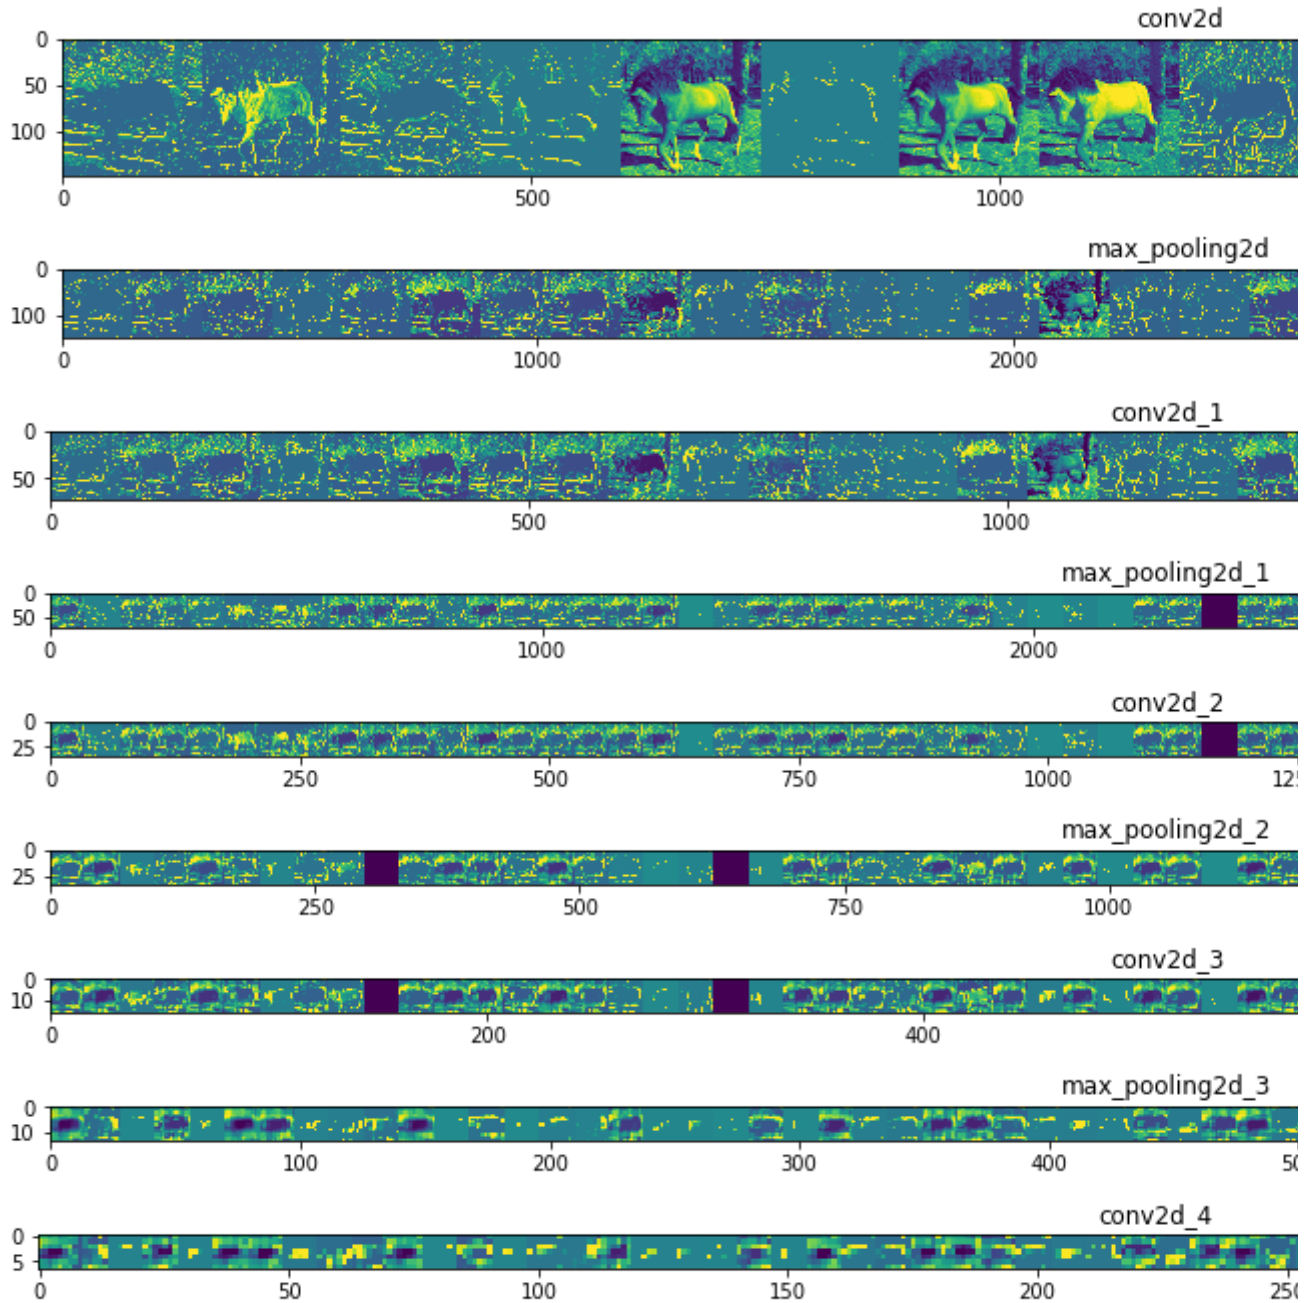


```
36 size = feature_map.shape[1]
37 # We will tile our images in this matrix
38 display_grid = np.zeros((size, size * n_features))
39 for i in range(n_features):
40     # Postprocess the feature to make it visually palatable
41     x = feature_map[0, :, :, i]
42     x -= x.mean()
43     x /= x.std()
44     x *= 64
45     x += 128
46     x = np.clip(x, 0, 255).astype('uint8')
47     # We'll tile each filter into this big horizontal grid
48     display_grid[:, i * size : (i + 1) * size] = x
49 # Display the grid
50 scale = 20. / n_features
51 plt.figure(figsize=(scale * n_features, scale))
52 plt.title(layer_name)
53 plt.grid(False)
54 plt.imshow(display_grid, aspect='auto', cmap='viridis')
```



Your session crashed for an unknown reason. [View runtime logs](#) ✕

/usr/local/lib/python3.6/dist-packages/ipykernel\_launcher.py:43: RuntimeWarni



As you can see we go from the raw pixels of the images to increasingly abstract and compact representations. Downstream layers start highlighting what the network pays attention to, and they show fewer and fewer non-zero values. This is called "sparsity." Representation sparsity is a key feature of deep learning.

These representations carry increasingly less information about the original pixels of the image, but they still contain enough information to classify the image. You can think of a convnet (or a deep network in general) as an information distillation process.

## ▼ Clean Up

Your session crashed for an unknown reason. [View runtime logs](#) ✕

Restart the kernel and free memory resources

```
1 import os, signal
2 os.kill(os.getpid(), signal.SIGKILL)
```

-- NORMAL --

Your session crashed for an unknown reason. [View runtime logs](#) ✕